Architecture

1. A spout connected to the Twitter streaming API, that pulls tweets and emits them to the parse bolt.
2. A parse-tweet-bolt, that parses the tweets emitted by the spout, and extracts individual words out of the received tweet text.
3. A Postgres database called **tcount**, with a table called **tweetwordcount**.
4. A count-bolt, that counts the number of words emitted by the tweet-parse bolt, and updates the total counts for each word in the Postgres table.

File structure

1. Create database and table: Extweetwordcount/psycopg-create.py
2. Spout connected to Twitter streaming: Extweetwordcount/src/spouts/tweets.py
3. Parse-tweet-bolt: Extweetwordcount/src/bolts/parse.py
4. Count-bolt: Extweetwordcount/src/bolts/wordcount.py
5. Topologies: Extweetwordcount/topologies/tweetwordcount.clj

File dependencies

1. RUN psycopg-create.py
2. RUN tweetwordcount.clj --> tweets.py --> parse.py --> wordcount.py --> tweetwordcount